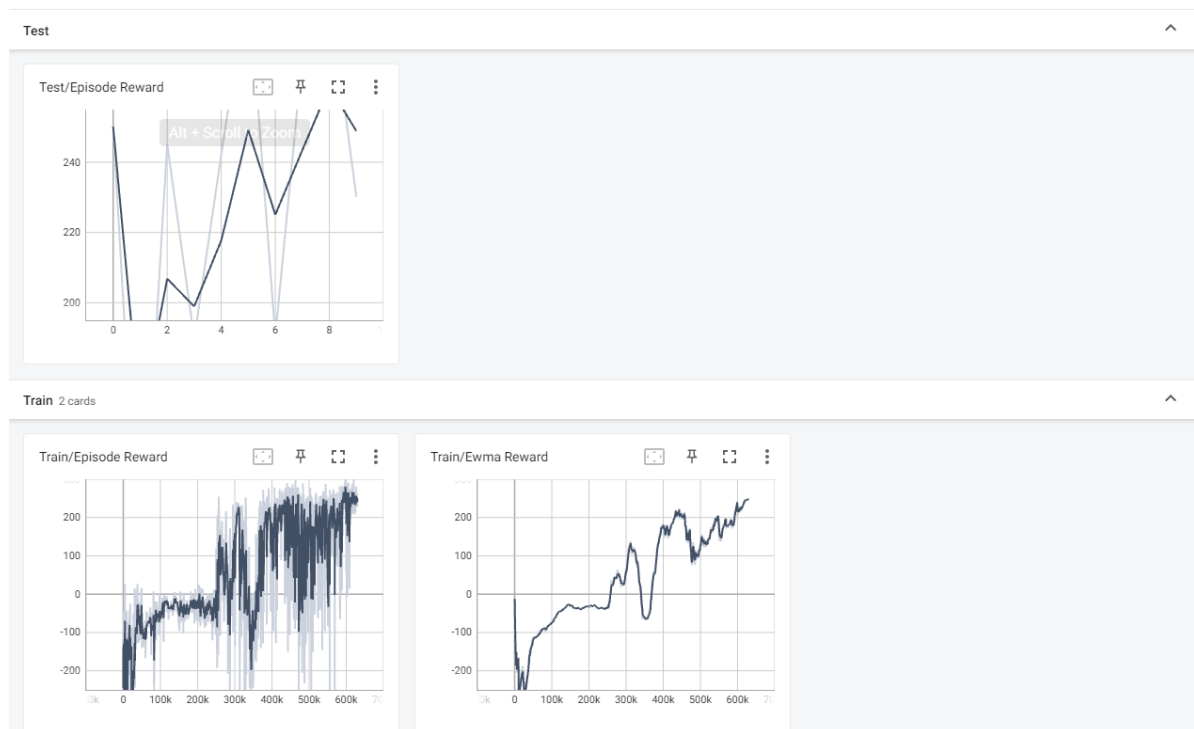
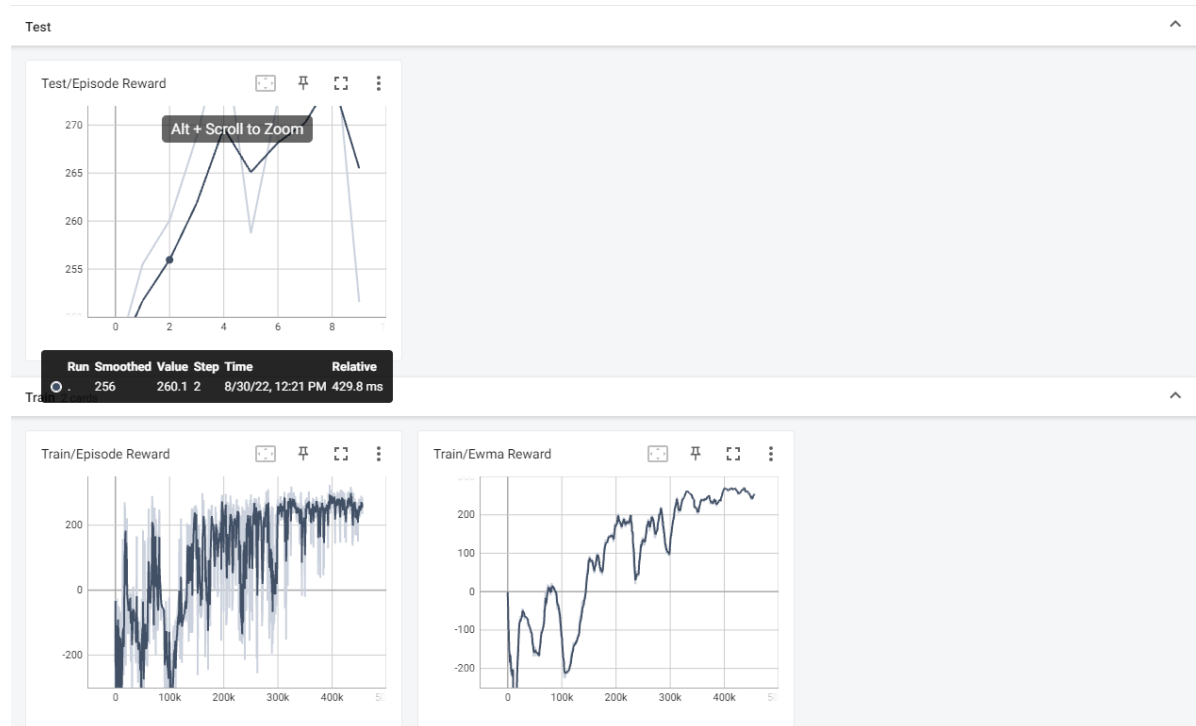


Lab6 - Deep Q-Network and Deep Deterministic Policy Gradient

- Student Info
 1. Student ID: 310555024
 2. Student Name: 林廷翰
- A tensorboard plot shows episode rewards of at least 800 training episodes in LunarLander-v2 (5%)



- A tensorboard plot shows episode rewards of at least 800 training episodes in LunarLanderContinuous-v2 (5%)



- Describe your major implementation of both algorithms in detail. (20%)

1. DQN

- Net

```

41 class Net(nn.Module):
42     def __init__(self, state_dim=8, action_dim=4, hidden_dim=32):
43         super().__init__()
44         # TODO
45         self.network = nn.Sequential(
46             nn.Linear(in_features=state_dim,
47                       out_features=hidden_dim),
48             nn.ReLU(inplace=True),
49             nn.Linear(in_features=hidden_dim,
50                       out_features=hidden_dim),
51             nn.ReLU(inplace=True),
52             nn.Linear(in_features=hidden_dim,
53                       out_features=action_dim)
54         )
55
56     def forward(self, x):
57         # TODO
58         return self.network(x)

```

在DQN中用3 layers的网络來預測 $Q(s, a)$ 的值，因為最後分類結果action有4種可 (No-op, Fire left engine, Fire main engine, Fire right engine)，所以最後一層 $\text{dim}=4$ 。

- select_action

```

81 | # epsilon-greedy based on behavior network.
82 | def select_action(self, state, epsilon, action_space):
83 |     # TODO
84 |     if random.random() > epsilon:
85 |         state = torch.from_numpy(state).float().unsqueeze(0).to(self.device)
86 |         self._behavior_net.eval()
87 |         with torch.no_grad():
88 |             action_values = self._behavior_net(state)
89 |             self._behavior_net.train()
90 |             return np.argmax(action_values.cpu().data.numpy())
91 |     else:
92 |         return random.choice(np.arange(action_space.n))

```

select action在每輪episode中選擇最大 $Q(s, a_i)$ 的 a_i 或有一定機率選擇。

- _update_behavior_network

```

106 | # Update behavior network.
107 | def _update_behavior_network(self, gamma):
108 |     # sample a minibatch of transitions
109 |     state, action, reward, next_state, done = self._memory.sample(
110 |         self.batch_size, self.device)
111 |
112 |     # TODO
113 |     q_value = self._behavior_net(state).gather(1, action.long())
114 |     with torch.no_grad():
115 |         q_next = self._target_net(next_state).detach().max(1)[0].unsqueeze(1)
116 |         q_target = reward + (gamma * q_next * (1 - done))
117 |         loss = nn.MSELoss()(q_value, q_target)
118 |
119 |     # Optimize
120 |     self._optimizer.zero_grad()
121 |     loss.backward()
122 |     nn.utils.clip_grad_norm_(self._behavior_net.parameters(), 5)
123 |     self._optimizer.step()

```

update network是由replay memory中sampling一些遊戲過程來做TD learning，再用 $Q(s, a)$ 與 $r + \gamma \max_a' Q(s', a')$ 的差做 MSELoss。

- _update_target_network

```

125 | # update target network by copying from behavior network.
126 | def _update_target_network(self):
127 |     # TODO
128 |     self._target_net.load_state_dict(self._behavior_net.state_dict())

```

最後每隔一段時間就用behavior network取代target network。

2. DDPG

- ActorNet

```

50 class ActorNet(nn.Module):
51     def __init__(self, state_dim=8, action_dim=2, hidden_dim=(400, 300)):
52         super().__init__()
53         # TODO
54         h1, h2 = hidden_dim
55         self.network = nn.Sequential(
56             nn.Linear(state_dim, h1),
57             nn.ReLU(inplace=True),
58             nn.Linear(h1, h2),
59             nn.ReLU(inplace=True),
60             nn.Linear(h2, action_dim),
61             nn.Tanh()
62         )
63
64     def forward(self, x):
65         # TODO
66         return self.network(x)

```

actor network可以根據當前的state決定下個要執行的action, 因為action有2個(main engine, left-right-engine), 所以最後輸出 dim=2。

- CriticNet

```

69 class CriticNet(nn.Module):
70     def __init__(self, state_dim=8, action_dim=2, hidden_dim=(400, 300)):
71         super().__init__()
72         h1, h2 = hidden_dim
73         self.critic_head = nn.Sequential(
74             nn.Linear(state_dim + action_dim, h1),
75             nn.ReLU(),
76         )
77         self.critic = nn.Sequential(
78             nn.Linear(h1, h2),
79             nn.ReLU(),
80             nn.Linear(h2, 1),
81         )
82
83     def forward(self, x, action):
84         x = self.critic_head(torch.cat([x, action], dim=1))
85         return self.critic(x)

```

critic network可以預估Q(s, a), 由於輸出為scalar, 所以輸出 dim=1。

- select_action

```

115 # Select an action based on the behavior (actor) network and exploration noise.
116 def select_action(self, state, noise=True):
117     # TODO
118     state = torch.from_numpy(state).float().to(self.device)
119
120     self._actor_net.eval()
121     with torch.no_grad():
122         action = self._actor_net(state).cpu().data.numpy()
123     self._actor_net.train()
124
125     if noise:
126         action += self._action_noise.sample()
127
128     return action

```

由actor network 選擇action並加上noise。

- `_update_behavior_network`

```

155 # Update critic
156 # Critic loss
157 # TODO
158 q_value = critic_net(state, action)
159 with torch.no_grad():
160     a_next = target_actor_net(next_state)
161     q_next = target_critic_net(next_state, a_next)
162     q_target = reward + (gamma * q_next * (1 - done))
163 critic_loss = nn.MSELoss()(q_value, q_target)

```

在每輪episode裡更新behavior的actor network μ , critic network Q , target 的 actor network μ' , critic network Q' 。再利用target network輸出的 q_target 和behavior network輸出的 q_value 取得MSE loss。

```

171 # Update actor
172 # Actor loss
173 # TODO
174 action = actor_net(state)
175 actor_loss = -critic_net(state, action).mean()

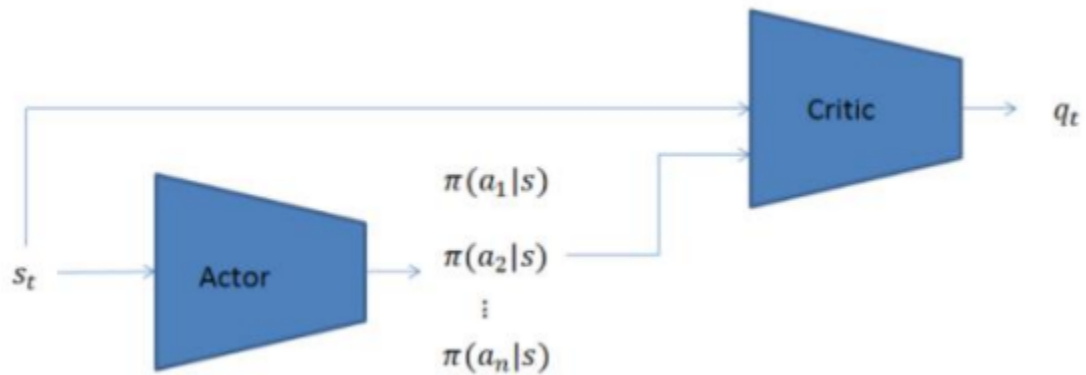
```

用behavior network的actor network μ 跟critic network Q 求得 $Q(s, a)$, 並且定義 loss function $E[-Q(a, \mu(s))]$ 更新 μ 使得輸出的 $Q(s, a)$ 越大越好, 並透過 bp 更新網路。

- Describe the differences between your implementation and algorithms. (10%)

在剛開始時會有一段時間不更新 network 中的參數而執行遊戲(隨機選action), 並把遊戲過程存到 replay memory 中, 且 DQN 部分中每隔 4 個iterations 才更新 behavior network。

- Describe your implementation and the gradient of actor updating. (10%)



```

171 # Update actor
172 # Actor loss
173 # TODO
174 action = actor_net(state)
175 actor_loss = -critic_net(state, action).mean()
176
177 # Optimize actor
178 actor_net.zero_grad()
179 critic_net.zero_grad()
180 actor_loss.backward()
181 actor_opt.step()
  
```

可以利用behavior network中的actor network μ 和critic network Q 求出 $Q(s, a)$ ，這時定義loss function為 $-Q(s, \mu(s))$ 來更新actor network μ ，bp時只更新actor，不更新critic，使得輸出 $Q(s, a)$ 越大越好。

- Describe your implementation and the gradient of critic updating. (10%)

```

155 # Update critic
156 # Critic loss
157 # TODO
158 q_value = critic_net(state, action)
159 with torch.no_grad():
160     a_next = target_actor_net(next_state)
161     q_next = target_critic_net(next_state, a_next)
162     q_target = reward + (gamma * q_next * (1 - done))
163 critic_loss = nn.MSELoss()(q_value, q_target)
164
165 # Optimize critic
166 actor_net.zero_grad()
167 critic_net.zero_grad()
168 critic_loss.backward()
169 critic_opt.step()
  
```

critic是利用target network輸出的 Q_{target} 和behavior network輸出的 $Q(s, a)$ 做MSE來更新 Q network。

- Explain the effects of the discount factor. (5%)

$$G_t = R_{t+1} + \lambda R_{t+2} + \dots = \sum_{k=0}^{\infty} \lambda^k R_{t+k+1}$$

λ 即為 discount factor, 意即越後面的 reward 的影響越小, 當前 reward 的影響最大。

- Explain the benefits of epsilon-greedy in comparison to greedy action selection. (5%)

用 epsilon-greedy 的方式因為偶爾會選擇其他隨機 action, 可以避免陷入局部最佳解的困境。

- Explain the necessity of the target network. (5%)

有 target network 和 behavior network 搭配使 training 更穩定, 因為 target network 每隔一段時間才會更新。

- Explain the effect of replay buffer size in case of too large or too small. (5%)

當 replay buffer size 越大, training 可以更穩定, 但相對會降低速度。當它太小時, 會一直著重在最近幾次 episode 中, 造成 overfitting。

- Implement and experiment on Double-DQN

1. Overview

DDQN 和 DQN 在實際上差不多, 差別只在更新 behavior network 時是如何決定 q_{target} 的; DDQN 是用 $\max_i Q(s, a_i)$ 作為找 $Q'(s, a_i)$ 的 index, 而不是直接取 $\max Q'(s, a_i)$ 。

2. `_update_behavior_network`

```

111 # Update behavior network.
112 def _update_behavior_network(self, gamma):
113     # sample a minibatch of transitions
114     state, action, reward, next_state, done = self._memory.sample(
115         self.batch_size, self.device)
116
117     # TODO
118     q_value = self._behavior_net(state).gather(1, action.long())
119     with torch.no_grad():
120         q_argmax = self._behavior_net(next_state).detach().max(1)[1].unsqueeze(1)
121         q_next = self._target_net(next_state).detach().gather(1, q_argmax)
122         q_target = reward + (gamma * q_next * (1 - done))
123
124     loss = nn.MSELoss()(q_value, q_target)
125
126     # Optimize
127     self._optimizer.zero_grad()
128     loss.backward()
129     nn.utils.clip_grad_norm_(self._behavior_net.parameters(), 5)
130     self._optimizer.step()

```

3. Average reward

```

Step: 606202 Episode: 1192 Length: 620 Total reward: -111.91 Ewma reward: 180.73 Epsilon: 0.010
Step: 606717 Episode: 1193 Length: 515 Total reward: 217.59 Ewma reward: 182.57 Epsilon: 0.010
Step: 606886 Episode: 1194 Length: 169 Total reward: 236.63 Ewma reward: 185.28 Epsilon: 0.010
Step: 607047 Episode: 1195 Length: 161 Total reward: -54.24 Ewma reward: 173.30 Epsilon: 0.010
Step: 607237 Episode: 1196 Length: 190 Total reward: 259.72 Ewma reward: 177.62 Epsilon: 0.010
Step: 607416 Episode: 1197 Length: 179 Total reward: 16.31 Ewma reward: 169.56 Epsilon: 0.010
Step: 607561 Episode: 1198 Length: 145 Total reward: -31.04 Ewma reward: 159.53 Epsilon: 0.010
Step: 607835 Episode: 1199 Length: 274 Total reward: 273.74 Ewma reward: 165.24 Epsilon: 0.010
Start Testing
C:\Users\user\Documents\Course\DL\venv\lib\site-packages\gym\core.py:256: DeprecationWarning: WARN: Fun
deprecation(
Average Reward 118.67440647932713

Process finished with exit code 0

```

- [LunarLander-v2] Average reward of 10 testing episodes: Average ÷ 30 (10%)

```

Step: 628012 Episode: 1192 Length: 238 Total reward: 227.22 Ewma reward: 246.92 Epsilon: 0.010
Step: 628280 Episode: 1193 Length: 268 Total reward: 267.10 Ewma reward: 247.93 Epsilon: 0.010
Step: 628503 Episode: 1194 Length: 223 Total reward: 245.79 Ewma reward: 247.82 Epsilon: 0.010
Step: 628807 Episode: 1195 Length: 304 Total reward: 250.61 Ewma reward: 247.96 Epsilon: 0.010
Step: 629108 Episode: 1196 Length: 301 Total reward: 247.60 Ewma reward: 247.94 Epsilon: 0.010
Step: 629378 Episode: 1197 Length: 270 Total reward: 253.58 Ewma reward: 248.22 Epsilon: 0.010
Step: 629645 Episode: 1198 Length: 267 Total reward: 247.12 Ewma reward: 248.17 Epsilon: 0.010
Step: 629869 Episode: 1199 Length: 224 Total reward: 228.57 Ewma reward: 247.19 Epsilon: 0.010
Start Testing
C:\Users\user\Documents\Course\DL\venv\lib\site-packages\gym\core.py:256: DeprecationWarning: WARN: Fun
deprecation(
Average Reward 231.5979033335987

Process finished with exit code 0

```


- [LunarLanderContinuous-v2] Average reward of 10 testing episodes: Average \div 30 (10%)

```

Step: 454162   Episode: 1192   Length: 199   Total reward: 252.10   Ewma reward: 252.16
Step: 454317   Episode: 1193   Length: 155   Total reward: 277.31   Ewma reward: 253.42
Step: 454592   Episode: 1194   Length: 275   Total reward: 235.01   Ewma reward: 252.50
Step: 454827   Episode: 1195   Length: 235   Total reward: 264.40   Ewma reward: 253.09
Step: 455012   Episode: 1196   Length: 185   Total reward: 270.73   Ewma reward: 253.97
Step: 455205   Episode: 1197   Length: 193   Total reward: 242.55   Ewma reward: 253.40
Step: 455428   Episode: 1198   Length: 223   Total reward: 284.38   Ewma reward: 254.95
Step: 455622   Episode: 1199   Length: 194   Total reward: 233.99   Ewma reward: 253.90
Start Testing
C:\Users\user\Documents\Course\DL\venv\lib\site-packages\gym\core.py:256: DeprecationWarning:
    deprecation(
Average Reward 264.7390025096876

Process finished with exit code 0

```

- Reference
 1. https://www.youtube.com/watch?v=mv0kfieln3s&ab_channel=NextDayVideo